

HARSH RAJ

harsh777111raj@gmail.com

harshraj172.github.io

[Google Scholar](#) [Github](#) [LinkedIn](#)

EDUCATION

Northeastern University

May 2026 (expected)

M.S. in Computer Science

GPA: 3.6/4.0

Related courses: Foundations of AI, Algorithms, Human-Computer Interaction

Delhi Technological University (DTU)

August 2023

B.Tech. in Engineering Physics

GPA: 3.4/4.0

Related courses: Machine Learning, Algorithmic Game Theory, Computer Vision.

RESEARCH INTERESTS

I am interested in large-scale data processing systems and training large models at scale.

EXPERIENCE

Research Assistant, Forschungszentrum Jülich

July 2025 – Present

Remote

Cologne, Germany

- Contributing to European Union's Open Foundational Model projects, including [MINERVA](#), [OpenEuroLLM](#), and [ELLIOT](#).

Co-Founder, Ontocord AI

Oct 2024 – Present

Remote

California, US

- Released the [MixtureVitae](#) dataset - the strongest permissibly licensed pretraining dataset to date.

Research Fellow, AI Risk and Vulnerability Alliance (ARVA)

Oct 2024 – June 2025

Remote

Washington, US

- Contributed multiple bias benchmarks to the UK AISI [inspect-ai](#) framework.
- Integrated the inspect-ai framework with the [AVID database](#).

Applied Scientist, Vijil AI

Feb 2024 – Oct 2024

Remote

California, US

- Developed [vijil-fuzzer](#), an LLM red-teaming framework for enterprise-grade customers. At the launch it got 80% attack rate against gpt-4o.
- Integrated open-source benchmarks into the [vijil-evaluation](#) service.
- Contributed to [vijil-dome](#), integrating guardrails into tools like Langchain with < 10ms of latency impact.

ML Engineer, Yield Protocol

Sept 2022 – Jan 2024

Remote

Chicago, US

- Led data synthesis to improve LLM agents and built [synchaev](#), a framework for agent data generation.
- Developed [mandrill](#) for finetuning and testing agents. Supervised 10 research interns at UIUC's Disruption Lab as part of the project.
- Contributed to the development of [Cacti](#), a Web3 transaction chatbot.

PUBLICATIONS

- **Harsh Raj**, Vipul Gupta, Domenic Rosati, Subhabrata Majumdar. *Improving Consistency in Large Language Models through Chain of Guidance* (TMLR'2025)
- Domenic Rosati, Giles Edkins, **Harsh Raj**, David Atanasov, Kai Williams, Subhabrata Majumdar, Janarthanan Rajendran, Frank Rudzicz, Hassan Sajjad. *Defences against Reverse Preference Attacks* (AI Security Workshop, AAAI'25)
- Laura Fee Nern, **Harsh Raj**, Maurice Georgi, Yash Sharma. *On Transfer of Adversarial Robustness from Pretraining to Downstream Tasks* (NeurIPS'23)
- **Harsh Raj**, Domenic Rosati, Subhabrata Majumdar. *Measuring Reliability of Large Language Models through Semantic Consistency* (Best Paper Award, ML Safety Workshop, NeurIPS'22)
- Sinchani Chakraborty, **Harsh Raj**, Srishti Gureja, Tanmay Jain, Atif Hassan, Sayantan Basu. *Evaluating the Robustness of Biomedical Concept Normalization* (Transfer Learning Workshop, NeurIPS'22)
- **Harsh Raj**, Ashutosh Pandey, Shaurya Kumar, Kavinder Singh, Nihal Kumar, Anil Singh Parihar. *Decoding Percepts in Vision Language Navigation: Is it about better features or more data?* ACM TIST
- Manu Joseph, **Harsh Raj**. *GANDALF: Gated Adaptive Network for Deep Automated Learning of Features* (under review, TMLR)
- Manu Joseph, **Harsh Raj**, Abhinav Yadav, Aaryamann Sharma. *AskYourDB: An end-to-end system for querying and visualizing relational databases using natural language*
- **Harsh Raj**, Aakansha Gupta, Rahul Katarya. *Extract It! Product Category Extraction by Transfer Learning* (CICT'22)

PROJECTS AND LIBRARIES

- [terminal-bench](#) (600 GitHub ★): Language Agent Sandbox Evaluation
Co-authored with Mike Meryll, Alex Shaw, Nicholas Carlini
- [inspect-ai](#) (1.3k GitHub ★): LLM Evaluation Suite.
Contributor
- [garak](#) (5k GitHub ★): LLM Red-Teaming Suite
Contributor

AWARDS, GRANTS AND HONORS

[SWISS AI Grant](#): Led the grant to acquire infrastructure for the project "From training a model to raising a model". The National Supercomputing Center (CSCS) granted access to 200k GPU hours on the Alps supercomputer and 61k CHF in monetary funds. 2025

[EuroHPC](#): As a co-founder at OntocordAI, acquired 4M GPU hours on the LEONARDO supercomputer for large-scale language model training. 2024

[Vision and Language Navigation](#): Secured 3rd position on the R2R benchmark 2024

[ML Safety Challenge](#): Best Paper Award with a \$5000 prize 2022

[Codeforces ML Round](#): Rank 7 in the Raif ML Round 1 2021