

HARSH RAJ

harsh777111raj@gmail.com

[harshraj172.github.io](https://github.com/harshraj172)

Google Scholar $\diamond$ Github $\diamond$ LinkedIn

EDUCATION

Northeastern University

M.S. in Computer Science

GPA: 3.6/4.0

May 2026 (*expected*)

Delhi Technological University (DTU)

B.Tech. in Engineering Physics

GPA: 3.4/4.0

August 2023

EXPERIENCE

Research Assistant, Forschungszentrum Jülich

Remote

July 2025 – Present

Cologne, Germany

- Contributing to European Union's Open Foundational Model projects, including [MINERVA](#), [OpenEuroLLM](#), and [ELLIOT](#).

Co-Founder, Ontocord AI

Remote

Oct 2024 – Present

California, US

- Released the [MixtureVitae](#) dataset - the strongest permissibly licensed pretraining dataset to date.

Applied Scientist, Vijil AI

Remote

Aug 2023 – Oct 2024

California, US

- Developed [vijil-fuzzer](#), an LLM red-teaming framework for enterprise-grade customers. At the launch it got 80% attack rate against gpt-4o.
- Integrated open-source benchmarks into the [vijil-evaluation](#) service.

SELECTED PUBLICATIONS

- **Harsh Raj**, Vipul Gupta, Domenic Rosati, Subhabrata Majumdar. *Improving Consistency in Large Language Models through Chain of Guidance* (TMLR'2025)
- Domenic Rosati, Giles Eddins, **Harsh Raj**, David Atanasov, Kai Williams, Subhabrata Majumdar, Janarthanan Rajendran, Frank Rudzicz, Hassan Sajjad. *Defences against Reverse Preference Attacks* (AI Security Workshop, AAAI'25)
- Laura Fee Nern, **Harsh Raj**, Maurice Georgi, Yash Sharma. *On Transfer of Adversarial Robustness from Pretraining to Downstream Tasks* (NeurIPS'23)
- **Harsh Raj**, Domenic Rosati, Subhabrata Majumdar. *Measuring Reliability of Large Language Models through Semantic Consistency* (Best Paper Award, ML Safety Workshop, NeurIPS'22)

PROJECTS AND LIBRARIES

- [terminal-bench](#) (600 GitHub ★): Language Agent Sandbox Evaluation
Co-authored with Mike Meryll, Alex Shaw, Nicholas Carlini
- [inspect-ai](#) (1.3k GitHub ★): LLM Evaluation Suite.
Contributor
- [garak](#) (5k GitHub ★): LLM Red-Teaming Suite
Contributor

AWARDS, GRANTS AND HONORS

[SWISS AI Grant](#): Awarded 200k GPU hours and CHF 61k in funding.

2025

[Vision and Language Navigation](#): Secured 3rd position on the R2R benchmark

2024

[ML Safety Challenge](#): Best Paper Award with a \$5000 prize

2022

[Codeforces ML Round](#): Rank 7 in the Raif ML Round 1

2021